

ISE - 364/464: Introduction to Machine Learning

Homework Assignment 7

The goal of this assignment is to provide a series of problems that solidify knowledge in specific topics in deep learning and unsupervised learning. Specifically, multi-layer perceptrons, K-means clustering, Gaussian mixture models, as well as connecting ridge regression to Bayesian linear regression.

Grading: This assignment is due on Coursesite by E.O.D. 12/13/2024. All problems are worth the same number of points. If a problem has multiple parts, each of those parts will be worth equal amounts and will sum to the total number of points of the original problem (Example: If each problem is worth a single point, and problem 1 has 4 parts, each part will be worth 1/4th of a point). ISE - 364 students are only required to answer problems 1 through 4; however, you are allowed to answer the 5th graduate-level question (if done so correctly, you will receive extra credit in the amount that the 5th problem will be worth for the ISE - 464 students). ISE - 464 students are required to answer all 5 problems.

Submitting: Only electronic submissions on Coursesite are accepted.

1 Problems

1. (MLP for Multi-Target Regression) Consider a multi-layer perception model $h_\theta(x)$ with some general number of layers $L > 1$ and predefined activation functions $g^{[\ell]}$, where the parameters $\theta := \{W, b\}$ denote the collection of weight matrices and bias vectors for each layer of the network, i.e., $W := \{W^{[\ell]}\}_{\ell=1}^L$ and $b := \{b^{[\ell]}\}_{\ell=1}^L$ where $W^{[\ell]} \in \mathbb{R}^{n^{[\ell]} \times n^{[\ell-1]}}$ and $b^{[\ell]} \in \mathbb{R}^{n^{[\ell]}}$ for all $\ell \in \{1, 2, \dots, L\}$. Further, this neural network is utilized to predict a multi-target continuous variable, i.e., this is a regression problem where the target variable is a vector $y \in \mathbb{R}^K$ where K is the number of targets. The squared error loss function is suitable for training this network, and, for a single feature-target pair (x, y) , is given by

$$J(\theta) := \frac{1}{2} \|h_\theta(x) - y\|_2^2.$$

- a) **(0.8 points)** Assuming that the activation function of the final layer is simply the identity function, i.e., $g^{[L]}(x) = I(x) = x$, derive expressions for the four following partial derivatives:

$$\frac{\partial J}{\partial a^{[L]}}, \quad \frac{\partial J}{\partial z^{[L]}}, \quad \frac{\partial J}{\partial W^{[L]}}, \quad \frac{\partial J}{\partial b^{[L]}}.$$

Also, write the dimensions of the resulting partial derivatives.

- b) **(0.2 points)** Suppose that we did not want to use the identity function in the final layer. Could one utilize the softmax function for the multi-target regression problem? Justify your answer.

2. **(Visualizing Feature Maps of an MLP)** Consider a general MLP $h_\theta(x)$ (i.e., a general number of layers $L > 1$ and activation functions g) where the parameters $\theta := \{W, b\}$ denote the collection of weight matrices and bias vectors for each layer of the network, i.e., $W := \{W^{[\ell]}\}_{\ell=1}^L$ and $b := \{b^{[\ell]}\}_{\ell=1}^L$ where $W^{[\ell]} \in \mathbb{R}^{n^{[\ell]} \times n^{[\ell-1]}}$ and $b^{[\ell]} \in \mathbb{R}^{n^{[\ell]}}$ for all $\ell \in \{1, 2, \dots, L\}$. Derive the partial derivative $\frac{\partial z_c^{[L]}}{\partial a^{[0]}}$, where $z_c^{[L]} = (W^{[L]})_c^\top a^{[L-1]} + b^{[L]}$, $(W^{[L]})_c^\top$ is the c -th row of the weight matrix $W^{[L]}$, and where $a^{[0]}$ is the input vector.
Hint: To do this, you will need to first derive an expression for $\frac{\partial z_c^{[L]}}{\partial z^{[L-1]}}$ and then a general expression for $\frac{\partial z_c^{[L]}}{\partial z^{[\ell]}}$ for all $1 \leq \ell \leq L$. Lastly, you should be able to define $\frac{\partial z_c^{[L]}}{\partial a^{[0]}}$ in terms of $\frac{\partial z_c^{[L]}}{\partial z^{[1]}}$.)

3. **(K-Means Clustering)** Recall that for some dataset $\mathcal{D} = \{x^{(i)}\}_{i=1}^m$, training a K-means clustering model with $K \in \mathbb{N}$ clusters amounts to solving the problem

$$\min_{r, \mu} J(r, \mu) := \frac{1}{2} \sum_{i=1}^m \sum_{k=1}^K r_{i,k} \|x^{(i)} - \mu_k\|_2^2,$$

where $r_{i,k} \in \{0, 1\}$ are the cluster assignment variables (with values of 1 if the datapoint $x^{(i)}$ is assigned to cluster k , and 0 otherwise) and $\mu_k \in \mathbb{R}^n$ are the cluster centroids, for all $i \in \{1, 2, \dots, m\}$ and $k \in \{1, 2, \dots, K\}$.

- a) Assume that each cluster has at least one datapoint assigned to it (i.e., $r_{i,k} = 1$ for at least one datapoint $x^{(i)}$ for all $k \in \{1, 2, \dots, K\}$). Prove that the Hessian of the K-means clustering loss function J is positive definite in the centroid variables μ_k when the cluster assignment variables $r_{i,k}$ are held constant for all $i \in \{1, 2, \dots, m\}$ and $k \in \{1, 2, \dots, K\}$.
- b) Derive the optimal solution for the centroids μ_k^* when the cluster assignment variables $r_{i,k}$ are held constant for all $i \in \{1, 2, \dots, m\}$ and $k \in \{1, 2, \dots, K\}$.
4. **(Gaussian Mixture Models)** Consider a Gaussian Mixture Model with $K \in \mathbb{N}$ clusters. Then, the incomplete-data log-likelihood function for a single observed datapoint $x \in \mathbb{R}^n$ corresponding to this model is given by

$$\ell(\theta; x) := \log \left\{ \sum_{k=1}^K \phi_k \mathcal{N}(x | \mu_k, \Sigma_k) \right\},$$

where the parameters are $\theta := \{\phi_1, \phi_2, \dots, \phi_K, \mu_1, \mu_2, \dots, \mu_K, \Sigma_1, \Sigma_2, \dots, \Sigma_K\}$, where $\phi_k \in (0, 1)$ denotes the probability of the datapoint coming from the k -th Gaussian cluster, and where $\mu_k \in \mathbb{R}^n$ and $\Sigma_k \in \mathbb{R}^{n \times n}$ denote the mean and covariance matrix corresponding to the k -th Gaussian cluster, respectively, for all $k \in \{1, 2, \dots, K\}$. While training a GMM when using the EM algorithm, one will need to compute the partial derivative of this log-likelihood with respect to μ_k . Prove that the partial derivative is

$$\frac{\partial \ell}{\partial \mu_k} := \frac{\phi_k \mathcal{N}(x | \mu_k, \Sigma_k)}{\sum_{j=1}^K \phi_j \mathcal{N}(x | \mu_j, \Sigma_j)} \Sigma_k^{-1} (x - \mu_k).$$

5. (ISE-464 Graduate Students) (Bayesian and Ridge Regression) Bayesian Linear Regression is a way of modeling a linear regression model $h_\theta(x) = \theta^\top x$ (with θ and x in \mathbb{R}^n , and the first entry of x is set to 1 to indicate the intercept) where the parameters are modeled as a multivariate Gaussian with mean 0 and covariance matrix $\tau^2 I$ (i.e., $\theta \sim \mathcal{N}(0, \tau^2 I)$ for some $\tau \in \mathbb{R}$ and where $I \in \mathbb{R}^{n \times n}$ denotes the identity matrix). Further, the target variable $y^{(i)} \in \mathbb{R}$ is still modeled as a Gaussian random variable with mean $\theta^\top x^{(i)}$ and variance σ (i.e., $y^{(i)} \sim \mathcal{N}(\theta^\top x^{(i)}, \sigma)$). Utilizing Bayes Theorem, one can formulate the Maximum A Posterior (MAP) estimate of the Bayesian linear regression model, denoted as θ_{MAP}^* , that maximizes the posterior distribution, as the problem

$$\theta_{MAP}^* := \operatorname{argmax}_{\theta} \mathbb{P} \left(\theta \mid \left\{ \left(x^{(i)}, y^{(i)} \right) \right\}_{i=1}^m ; \tau, \sigma \right) = \operatorname{argmax}_{\theta} \mathbb{P}(\theta; \tau) \prod_{i=1}^m \mathbb{P} \left(y^{(i)} \mid x^{(i)}; \theta, \sigma \right),$$

where the prior distribution $\mathbb{P}(\theta; \tau)$ is given by

$$\mathbb{P}(\theta; \tau) = \frac{1}{(2\pi)^{n/2} |\tau^2 I|^{1/2}} e^{-\frac{1}{2} \theta^\top (\tau^2 I)^{-1} \theta},$$

and the posterior distribution $\mathbb{P}(y^{(i)} \mid x^{(i)}; \theta, \sigma)$ is given by

$$\mathbb{P} \left(y^{(i)} \mid x^{(i)}; \theta, \sigma \right) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{1}{2\sigma^2} (y^{(i)} - \theta^\top x^{(i)})^2}.$$

Prove that θ_{MAP}^* leads to the ℓ_2 -regularized linear regression problem (ridge regression), i.e., prove that

$$\theta_{MAP}^* = \operatorname{argmin}_{\theta} \frac{1}{2} \sum_{i=1}^m \left(\theta^\top x^{(i)} - y^{(i)} \right)^2 + \frac{\lambda}{2} \|\theta\|_2^2,$$

for some $\lambda > 0$.